

## ▼ Exploratory Data Analysis (EDA) and Data Visualization

### Module 4, Lab 2: Understanding Your Data

Exploratory Data Analysis (EDA) is one of the most critical steps in any machine learning project. Before building models, you need to understand your data thoroughly. This lab will teach you how to explore datasets, identify patterns, and create meaningful visualizations.

#### Learning Objectives

By the end of this lab, you will be able to:

- Load and examine datasets using pandas
- Identify data quality issues (missing values, duplicates, outliers)
- Calculate and interpret summary statistics
- Create effective visualizations using matplotlib and seaborn
- Draw insights from data exploration

#### Business Problem

We'll analyze a customer dataset to understand purchasing behavior and demographics. This type of analysis helps businesses make data-driven decisions about marketing, product development, and customer segmentation.

## ▼ Setup and Data Loading

```
# Install required packages
!pip install --upgrade pip
!pip install pandas numpy matplotlib seaborn plotly

Requirement already satisfied: pip in /usr/local/lib/python3.12/dist-packages
Collecting pip
  Downloading pip-25.3-py3-none-any.whl.metadata (4.7 kB)
  Downloading pip-25.3-py3-none-any.whl (1.8 MB)
  1.8/1.8

Installing collected packages: pip
  Attempting uninstall: pip
    Found existing installation: pip 24.1.2
    Uninstalling pip-24.1.2:
      Successfully uninstalled pip-24.1.2
  Successfully installed pip-25.3
Requirement already satisfied: pandas in /usr/local/lib/python3.12/dist-packages
Requirement already satisfied: numpy in /usr/local/lib/python3.12/dist-packages
```

```
Requirement already satisfied: matplotlib in /usr/local/lib/python3.12/dist-p  
Requirement already satisfied: seaborn in /usr/local/lib/python3.12/dist-pac  
Requirement already satisfied: plotly in /usr/local/lib/python3.12/dist-packa  
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/pytho  
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.12/dist  
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.12/di  
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.12/  
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.12/dist  
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.12  
Requirement already satisfied: kiwisolver>=1.3.1 in /usr/local/lib/python3.12  
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.12/d  
Requirement already satisfied: pillow>=8 in /usr/local/lib/python3.12/dist-pa  
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.12/  
Requirement already satisfied: tenacity>=6.2.0 in /usr/local/lib/python3.12/d  
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.12/dist-pac
```

```
# Import necessary libraries  
import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt  
import seaborn as sns  
import warnings  
warnings.filterwarnings('ignore')  
  
# Set plotting style  
plt.style.use('seaborn-v0_8')  
sns.set_palette("husl")  
%matplotlib inline  
  
# Set display options  
pd.set_option('display.max_columns', None)  
pd.set_option('display.width', None)  
  
print("Libraries imported successfully!")
```

Libraries imported successfully!

## >Loading the Dataset

We'll create a realistic customer dataset for our analysis.

```
# Create a synthetic customer dataset  
np.random.seed(42)  
n_customers = 1000  
  
# Generate customer data  
customer_data = {  
    'customer_id': range(1, n_customers + 1),  
    'age': np.random.normal(40, 15, n_customers).astype(int),
```

```

'gender': np.random.choice(['Male', 'Female', 'Other'], n_customers, p=[

'income': np.random.lognormal(10.5, 0.5, n_customers),

'education': np.random.choice(['High School', 'Bachelor', 'Master', 'PhD'],
                                n_customers, p=[0.3, 0.4, 0.25, 0.05]),

'city_tier': np.random.choice(['Tier 1', 'Tier 2', 'Tier 3'],
                                n_customers, p=[0.3, 0.4, 0.3]),

'years_as_customer': np.random.exponential(3, n_customers),
'total_purchases': np.random.poisson(12, n_customers),
'avg_order_value': np.random.gamma(2, 50, n_customers),
'satisfaction_score': np.random.normal(7.5, 1.5, n_customers)

}

# Create DataFrame
df = pd.DataFrame(customer_data)

# Add some realistic constraints
df['age'] = np.clip(df['age'], 18, 80)
df['income'] = np.clip(df['income'], 20000, 200000)
df['years_as_customer'] = np.clip(df['years_as_customer'], 0, 15)
df['satisfaction_score'] = np.clip(df['satisfaction_score'], 1, 10)

# Calculate total spending
df['total_spending'] = df['total_purchases'] * df['avg_order_value']

# Introduce some missing values (realistic scenario)
missing_indices = np.random.choice(df.index, size=50, replace=False)
df.loc[missing_indices, 'satisfaction_score'] = np.nan

# Add some duplicates (data quality issue)
duplicate_rows = df.sample(5).copy()
df = pd.concat([df, duplicate_rows], ignore_index=True)

print(f"Dataset created with {len(df)} customers")
print(f"Dataset shape: {df.shape}")

```

Dataset created with 1005 customers  
 Dataset shape: (1005, 11)

## Step 1: Initial Data Exploration

Let's start by getting familiar with our dataset.

```

# Display first few rows
print("First 5 rows of the dataset:")
display(df.head())

print("\nLast 5 rows of the dataset:")
display(df.tail())

```

First 5 rows of the dataset:

	customer_id	age	gender	income	education	city_tier	years_as_cust
0	1	47	Male	31113.449406	High School	Tier 2	0.4
1	2	37	Male	24932.359005	Bachelor	Tier 2	6.28
2	3	49	Female	42599.051008	Bachelor	Tier 1	2.25
3	4	62	Female	70985.096321	High School	Tier 1	1.8
4	5	36	Male	20000.000000	High School	Tier 1	0.33

Last 5 rows of the dataset:

	customer_id	age	gender	income	education	city_tier	years_as_c
1000	935	46	Female	96478.813777	Master	Tier 2	
1001	644	18	Male	30728.977611	Bachelor	Tier 3	
1002	810	26	Male	20000.000000	Bachelor	Tier 2	

```
# Get basic information about the dataset
print("Dataset Info:")
print(df.info())

print("\nDataset Shape:")
print(f"Rows: {df.shape[0]}, Columns: {df.shape[1]}")

print("\nColumn Names:")
print(df.columns.tolist())
```

```
Dataset Info:
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1005 entries, 0 to 1004
Data columns (total 11 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   customer_id      1005 non-null   int64  
 1   age              1005 non-null   int64  
 2   gender           1005 non-null   object  
 3   income           1005 non-null   float64
 4   education        1005 non-null   object  
 5   city_tier        1005 non-null   object  
 6   years_as_customer 1005 non-null   float64
 7   total_purchases  1005 non-null   int64  
 8   avg_order_value  1005 non-null   float64
 9   satisfaction_score 954 non-null   float64
 10  total_spending   1005 non-null   float64
dtypes: float64(5), int64(3), object(3)
```

```
memory usage: 86.5+ KB
None
```

```
Dataset Shape:
Rows: 1005, Columns: 11
```

```
Column Names:
['customer_id', 'age', 'gender', 'income', 'education', 'city_tier', 'years_a
```

```
# Check data types
print("Data Types:")
print(df.dtypes)

print("\nNumerical Columns:")
numerical_cols = df.select_dtypes(include=[np.number]).columns.tolist()
print(numerical_cols)

print("\nCategorical Columns:")
categorical_cols = df.select_dtypes(include=['object']).columns.tolist()
print(categorical_cols)
```

```
Data Types:
customer_id          int64
age                  int64
gender               object
income               float64
education            object
city_tier            object
years_as_customer    float64
total_purchases      int64
avg_order_value     float64
satisfaction_score  float64
total_spending       float64
dtype: object
```

```
Numerical Columns:
['customer_id', 'age', 'income', 'years_as_customer', 'total_purchases', 'avg
Categorical Columns:
['gender', 'education', 'city_tier']
```

## ▼ Step 2: Data Quality Assessment

Before analyzing the data, we need to identify and understand data quality issues.

```
# Check for missing values
print("Missing Values:")
missing_data = df.isnull().sum()
missing_percent = (missing_data / len(df)) * 100
```

```
missing_df = pd.DataFrame({
    'Missing Count': missing_data,
    'Missing Percentage': missing_percent
})
missing_df = missing_df[missing_df['Missing Count'] > 0].sort_values('Missing Count', ascending=False)
print(missing_df)
```

Missing Values:

	Missing Count	Missing Percentage
satisfaction_score	51	5.074627

```
# Check for duplicate rows
print(f"Total rows: {len(df)}")
print(f"Unique rows: {len(df.drop_duplicates())}")
print(f"Duplicate rows: {len(df) - len(df.drop_duplicates())}")

if len(df) != len(df.drop_duplicates()):
    print("\nDuplicate rows found:")
    duplicates = df[df.duplicated(keep=False)]
    print(duplicates.sort_values('customer_id'))
```

Total rows: 1005

Unique rows: 1000

Duplicate rows: 5

Duplicate rows found:

	customer_id	age	gender	income	education	city_tier
21	22	36	Female	46833.896860	Bachelor	Tier 2
1004	22	36	Female	46833.896860	Bachelor	Tier 2
1003	286	18	Male	67377.429094	Bachelor	Tier 1
285	286	18	Male	67377.429094	Bachelor	Tier 1
643	644	18	Male	30728.977611	Bachelor	Tier 3
1001	644	18	Male	30728.977611	Bachelor	Tier 3
809	810	26	Male	20000.000000	Bachelor	Tier 2
1002	810	26	Male	20000.000000	Bachelor	Tier 2
934	935	46	Female	96478.813777	Master	Tier 2
1000	935	46	Female	96478.813777	Master	Tier 2

	years_as_customer	total_purchases	avg_order_value	satisfaction_score
21	3.994514	11	42.648150	8.604621
1004	3.994514	11	42.648150	8.604621
1003	1.771158	8	31.961654	7.293824
285	1.771158	8	31.961654	7.293824
643	1.151774	12	92.621642	NaN
1001	1.151774	12	92.621642	NaN
809	0.227681	11	42.907979	2.973530
1002	0.227681	11	42.907979	2.973530
934	3.218394	11	137.763089	9.208171
1000	3.218394	11	137.763089	9.208171

	total_spending
21	469.129645
1004	469.129645

1003	255.693232
285	255.693232
643	1111.459705
1001	1111.459705
809	471.987767
1002	471.987767
934	1515.393977
1000	1515.393977

```
# Check for outliers using IQR method
def detect_outliers(df, column):
    Q1 = df[column].quantile(0.25)
    Q3 = df[column].quantile(0.75)
    IQR = Q3 - Q1
    lower_bound = Q1 - 1.5 * IQR
    upper_bound = Q3 + 1.5 * IQR
    outliers = df[(df[column] < lower_bound) | (df[column] > upper_bound)]
    return outliers, lower_bound, upper_bound

print("Outlier Detection (using IQR method):")
for col in ['age', 'income', 'total_spending']:
    outliers, lower, upper = detect_outliers(df, col)
    print(f"\n{col}:")
    print(f"  Normal range: {lower:.2f} to {upper:.2f}")
    print(f"  Number of outliers: {len(outliers)}")
    print(f"  Percentage of outliers: {len(outliers)/len(df)*100:.2f}%")
```

Outlier Detection (using IQR method):

age:  
Normal range: 1.50 to 77.50  
Number of outliers: 6  
Percentage of outliers: 0.60%

income:  
Normal range: -10132.92 to 87157.18  
Number of outliers: 36  
Percentage of outliers: 3.58%

total\_spending:  
Normal range: -913.79 to 3050.00  
Number of outliers: 43  
Percentage of outliers: 4.28%

## ▼ Step 3: Summary Statistics

Let's calculate and interpret summary statistics for our numerical variables.

```
# Basic summary statistics
print("Summary Statistics for Numerical Variables:")
summary_stats = df.describe()
display(summary_stats.round(2))
```

Summary Statistics for Numerical Variables:

	customer_id	age	income	years_as_customer	total_purchases	avg_order_value
<b>count</b>	1005.00	1005.00	1005.00		1005.00	1005.00
<b>mean</b>	500.69	40.10	41604.00		3.01	12.00
<b>std</b>	289.10	13.89	20933.79		2.97	3.42
<b>min</b>	1.00	18.00	20000.00		0.01	3.00
<b>25%</b>	251.00	30.00	26350.87		0.82	10.00
<b>50%</b>	501.00	40.00	36443.98		2.08	12.00
<b>75%</b>	751.00	49.00	50673.39		4.27	14.00
<b>max</b>	1000.00	80.00	174359.45		15.00	23.00

```
# Additional statistics
print("Additional Statistics:")
additional_stats = pd.DataFrame({
    'Skewness': df[numerical_cols].skew(),
    'Kurtosis': df[numerical_cols].kurtosis(),
    'Variance': df[numerical_cols].var()
})
display(additional_stats.round(3))
```

Additional Statistics:

	Skewness	Kurtosis	Variance	
<b>customer_id</b>	-0.002	-1.202	8.357904e+04	
<b>age</b>	0.328	-0.391	1.928100e+02	
<b>income</b>	1.938	6.351	4.382238e+08	
<b>years_as_customer</b>	1.615	2.666	8.814000e+00	
<b>total_purchases</b>	0.328	-0.035	1.168500e+01	
<b>avg_order_value</b>	1.443	2.818	4.905801e+03	
<b>satisfaction_score</b>	-0.237	-0.138	1.961000e+00	
<b>total_spending</b>	1.606	3.718	7.867163e+05	

```
# Summary for categorical variables
print("Summary for Categorical Variables:")
```

```

for col in categorical_cols:
    print(f"\n{col}:")
    value_counts = df[col].value_counts()
    percentages = df[col].value_counts(normalize=True) * 100
    summary = pd.DataFrame({
        'Count': value_counts,
        'Percentage': percentages
    })
    print(summary.round(2))

```

Summary for Categorical Variables:

gender:

	Count	Percentage
gender		
Female	519	51.64
Male	465	46.27
Other	21	2.09

education:

	Count	Percentage
education		
Bachelor	434	43.18
High School	309	30.75
Master	220	21.89
PhD	42	4.18

city\_tier:

	Count	Percentage
city_tier		
Tier 2	410	40.80
Tier 3	301	29.95
Tier 1	294	29.25

## ▼ Step 4: Data Visualization

Now let's create visualizations to better understand our data patterns.

### ▼ 4.1 Distribution of Numerical Variables

```

# Create histograms for numerical variables
fig, axes = plt.subplots(2, 3, figsize=(18, 12))
axes = axes.ravel()


numerical_vars = ['age', 'income', 'years_as_customer', 'total_purchases', '']

for i, var in enumerate(numerical_vars):
    axes[i].hist(df[var].dropna(), bins=30, alpha=0.7, color='skyblue', edge
    axes[i].set_title(f'Distribution of {var.replace("_", " ")}.title()')

```

```
axes[i].set_xlabel(var.replace("_", " ").title())
axes[i].set_ylabel('Frequency')
axes[i].grid(True, alpha=0.3)


plt.tight_layout()
plt.show()
```



```
# Box plots to identify outliers
fig, axes = plt.subplots(2, 3, figsize=(18, 12))
axes = axes.ravel()

for i, var in enumerate(numerical_vars):
    axes[i].boxplot(df[var].dropna())
    axes[i].set_title(f'Box Plot of {var.replace("_", " ")}.title()')
    axes[i].set_ylabel(var.replace("_", " ").title())
    axes[i].grid(True, alpha=0.3)

plt.tight_layout()
plt.show()
```




## 4.2 Categorical Variable Analysis

```
# Bar plots for categorical variables
fig, axes = plt.subplots(1, 3, figsize=(18, 6))

for i, var in enumerate(categorical_cols):
    value_counts = df[var].value_counts()
    axes[i].bar(value_counts.index, value_counts.values, alpha=0.7)
    axes[i].set_title(f'Distribution of {var.replace("_", " ").title()}')
    axes[i].set_xlabel(var.replace("_", " ").title())
    axes[i].set_ylabel('Count')
    axes[i].tick_params(axis='x', rotation=45)
    axes[i].grid(True, alpha=0.3)


plt.tight_layout()
plt.show()
```



```
# Pie charts for categorical variables
fig, axes = plt.subplots(1, 3, figsize=(18, 6))

for i, var in enumerate(categorical_cols):
    value_counts = df[var].value_counts()
    axes[i].pie(value_counts.values, labels=value_counts.index, autopct='%1.
    axes[i].set_title(f'Proportion of {var.replace("_", " ").title()}')
```

```
plt.tight_layout()
plt.show()
```




## ⌄ 4.3 Correlation Analysis

```
# Correlation matrix
correlation_matrix = df[numerical_cols].corr()

plt.figure(figsize=(12, 10))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', center=0,
            square=True, linewidths=0.5, fmt='.2f')
plt.title('Correlation Matrix of Numerical Variables')
plt.tight_layout()
plt.show()

# Print strong correlations
print("Strong Correlations (|r| > 0.5):")
for i in range(len(correlation_matrix.columns)):
    for j in range(i+1, len(correlation_matrix.columns)):
        corr_value = correlation_matrix.iloc[i, j]
        if abs(corr_value) > 0.5:
            print(f"{correlation_matrix.columns[i]} vs {correlation_matrix.c
```





## 4.4 Relationship Analysis

```
# Scatter plots for key relationships
fig, axes = plt.subplots(2, 2, figsize=(15, 12))

# Income vs Total Spending
axes[0, 0].scatter(df['income'], df['total_spending'], alpha=0.6)
axes[0, 0].set_xlabel('Income')
axes[0, 0].set_ylabel('Total Spending')
axes[0, 0].set_title('Income vs Total Spending')
axes[0, 0].grid(True, alpha=0.3)


# Age vs Total Purchases
axes[0, 1].scatter(df['age'], df['total_purchases'], alpha=0.6)
axes[0, 1].set_xlabel('Age')
axes[0, 1].set_ylabel('Total Purchases')
axes[0, 1].set_title('Age vs Total Purchases')
axes[0, 1].grid(True, alpha=0.3)

# Years as Customer vs Satisfaction Score
axes[1, 0].scatter(df['years_as_customer'], df['satisfaction_score'], alpha=0.6)
axes[1, 0].set_xlabel('Years as Customer')
axes[1, 0].set_ylabel('Satisfaction Score')
axes[1, 0].set_title('Years as Customer vs Satisfaction Score')
axes[1, 0].grid(True, alpha=0.3)

# Average Order Value vs Total Purchases
axes[1, 1].scatter(df['avg_order_value'], df['total_purchases'], alpha=0.6)
axes[1, 1].set_xlabel('Average Order Value')
axes[1, 1].set_ylabel('Total Purchases')
axes[1, 1].set_title('Average Order Value vs Total Purchases')
axes[1, 1].grid(True, alpha=0.3)

plt.tight_layout()
plt.show()
```





## 4.5 Group Analysis


```
# Spending patterns by gender
plt.figure(figsize=(15, 5))

plt.subplot(1, 3, 1)
df.boxplot(column='total_spending', by='gender', ax=plt.gca())
plt.title('Total Spending by Gender')
plt.suptitle('') # Remove default title

plt.subplot(1, 3, 2)
df.boxplot(column='income', by='education', ax=plt.gca())
plt.title('Income by Education Level')
plt.suptitle('') # Remove default title
plt.xticks(rotation=45)

plt.subplot(1, 3, 3)
df.boxplot(column='satisfaction_score', by='city_tier', ax=plt.gca())
plt.title('Satisfaction Score by City Tier')
plt.suptitle('') # Remove default title

plt.tight_layout()
plt.show()
```



```
# Group statistics
print("Average Total Spending by Gender:")
gender_spending = df.groupby('gender')['total_spending'].agg(['mean', 'median'])
print(gender_spending)

print("\nAverage Income by Education Level:")
education_income = df.groupby('education')['income'].agg(['mean', 'median'])
print(education_income)

print("\nAverage Satisfaction Score by City Tier:")
city_satisfaction = df.groupby('city_tier')['satisfaction_score'].agg(['mean'])
print(city_satisfaction)
```

Average Total Spending by Gender:  
 mean median std

gender	mean	median	std
Female	1254.28	1034.91	897.43
Male	1123.97	900.81	868.60
Other	1286.33	1091.70	955.32

Average Income by Education Level:

education	mean	median	std
Bachelor	41635.96	36521.03	20196.63
High School	40884.57	36054.03	21005.57
Master	41657.57	36738.09	20194.18
PhD	46286.02	42204.91	30066.34

Average Satisfaction Score by City Tier:

city_tier	mean	median	std
Tier 1	7.41	7.45	1.43
Tier 2	7.41	7.43	1.36
Tier 3	7.56	7.42	1.42

## ▼ Step 5: Advanced Visualizations

```
# Pair plot for key numerical variables
key_vars = ['age', 'income', 'total_spending', 'satisfaction_score']
sns.pairplot(df[key_vars + ['gender']].dropna(), hue='gender', diag_kind='hi')
plt.suptitle('Pair Plot of Key Variables by Gender', y=1.02)
plt.show()
```



Pair Plot of Key Variables by Gender



```
# Create customer segments based on spending and purchases
df['spending_category'] = pd.cut(df['total_spending'],
                                  bins=[0, df['total_spending'].quantile(0.33),
                                         df['total_spending'].quantile(0.67),
                                         df['total_spending'].max()],
                                  labels=['Low Spender', 'Medium Spender', 'Hi'])

# Visualize segments
plt.figure(figsize=(15, 5))


plt.subplot(1, 3, 1)
segment_counts = df['spending_category'].value_counts()
plt.pie(segment_counts.values, labels=segment_counts.index, autopct='%.1f%%')
plt.title('Customer Segments by Spending')

plt.subplot(1, 3, 2)
sns.boxplot(data=df, x='spending_category', y='satisfaction_score')
plt.title('Satisfaction Score by Spending Category')
plt.xticks(rotation=45)

plt.subplot(1, 3, 3)
sns.boxplot(data=df, x='spending_category', y='years_as_customer')
plt.title('Customer Tenure by Spending Category')
plt.xticks(rotation=45)

plt.tight_layout()
plt.show()
```

age income total\_spending satisfaction\_score years\_as\_customer



## ▼ Step 6: Key Insights and Findings

Let's summarize our key findings from the EDA.

```
# Calculate key business metrics
print("==== KEY BUSINESS INSIGHTS ====")
print(f"\n📊 Dataset Overview:")
print(f"  • Total customers analyzed: {len(df)}")
print(f"  • Data quality: {((len(df) - df.isnull().sum().sum()) / (len(df))) * 100} %")

print(f"\n💰 Financial Metrics:")
print(f"  • Average customer income: ${df['income'].mean():,.0f}")
print(f"  • Average total spending: ${df['total_spending'].mean():,.0f}")
print(f"  • Average order value: ${df['avg_order_value'].mean():,.0f}")
print(f"  • Total revenue: ${df['total_spending'].sum():,.0f}")

print(f"\n👤 Customer Demographics:")
print(f"  • Average age: {df['age'].mean():.1f} years")
print(f"  • Gender distribution: {dict(df['gender'].value_counts())}")
print(f"  • Average customer tenure: {df['years_as_customer'].mean():.1f} years")

print(f"\n😊 Customer Satisfaction:")
```

```
print(f"  • Average satisfaction score: {df['satisfaction_score'].mean():.1f}")
print(f"  • Highly satisfied customers (>8): {len(df[df['satisfaction_score']

print("\n⌚ Customer Segments:")
segment_stats = df.groupby('spending_category').agg({
    'total_spending': 'mean',
    'satisfaction_score': 'mean',
    'years_as_customer': 'mean'
}).round(2)
for segment in segment_stats.index:
    count = len(df[df['spending_category'] == segment])
    print(f"  • {segment}: {count} customers ({count/len(df)*100:.1f}%)")
    print(f"    - Avg spending: ${segment_stats.loc[segment, 'total_spending']:.2f}")
    print(f"    - Avg satisfaction: {segment_stats.loc[segment, 'satisfaction_score']:.1f}")

==== KEY BUSINESS INSIGHTS ====
```

#### Dataset Overview:

- Total customers analyzed: 1,005
- Data quality: 7.9% complete

#### Financial Metrics:

- Average customer income: \$41,604
- Average total spending: \$1,195
- Average order value: \$101
- Total revenue: \$1,200,632

#### Customer Demographics:

- Average age: 40.1 years
- Gender distribution: {'Female': np.int64(519), 'Male': np.int64(465), 'Other': np.int64(21)}
- Average customer tenure: 3.0 years

#### Customer Satisfaction:

- Average satisfaction score: 7.5/10
- Highly satisfied customers (>8): 337/954 (35.3%)

#### Customer Segments:

- Low Spender: 332 customers (33.0%)
  - Avg spending: \$413
  - Avg satisfaction: 7.4/10
- Medium Spender: 341 customers (33.9%)
  - Avg spending: \$990
  - Avg satisfaction: 7.6/10
- High Spender: 332 customers (33.0%)
  - Avg spending: \$2,186
  - Avg satisfaction: 7.4/10

## ▼ Challenge: Your Turn to Explore!

Now it's your turn to practice EDA skills. Complete the following tasks:

## Challenge 1: Create a new visualization

Create a visualization that shows the relationship between education level and average order value. What insights can you draw?

```
# Your code here for Challenge 1
# Hint: Try using a bar plot or box plot
# Challenge 1: Relationship between education level and average order value

import matplotlib.pyplot as plt
import seaborn as sns

# Calculate average order value by education level
education_avg = df.groupby('education')['avg_order_value'].mean().reset_index()

# Plot bar chart
plt.figure(figsize=(8,5))
sns.barplot(x='education', y='avg_order_value', data=education_avg, palette=)
plt.title('Average Order Value by Education Level')
plt.xlabel('Education Level')
plt.ylabel('Average Order Value ($)')
plt.show()

# Insight:
# Customers with higher education levels tend to have a higher average order
# suggesting they may have greater purchasing power or prefer premium produc
```

